

EDA-01 Know Your Data

1. Necessary Imports

```
In [4]: import pandas as pd  
import numpy as np  
from matplotlib import pyplot as plt  
import seaborn as sns
```

a) What is Pandas in Python?

Pandas is an open source Python library which provides necessary support for implementing Series and dataframe data structures. A Pandas Series is a one dimensional array with homogenous data . The size of series is immutable but values are mutable. A dataframe is a data structure to hold two dimensional array with heterogeneous data columns. Pandas Dataframe is the ideal data structure for EDA and Pre-Processing Tasks. Pandas Dataframe Features:

1. Reading different file types like Excel and CSV
2. Default and user defined indexing
3. Easy data access based on indices
4. Easy handling of Missing Values
5. Easy Data Splitting & Merging , Concatenation and integration
6. Easy data sorting and aggregation
7. Easy and efficient handling of time series data

b) Numpy

Numpy stands for numerical python and provides support for efficient mathematical operations on array type data

c) Data Dictionary

A data dictionary is a data structure in python which stores data in the form of key-value pairs.

d) matplotlib and seaborn

Matplotlib and seaborn are python graphic libraries . Matplotlib is used for basic plotting while seaborn which is based on matplotlib is used for advance statistical visualizations.

2.Creating a Data Dictionary

```
In [5]: data={'Name':['Ali','Ahmed','Sara','Tom','Mike','Maria'],
'Gender':['Male','Male','Female','Male','Male','Female'],
'Age':[20,30, 15, 24, 23, 20],
'Height':[5.2,5.4,5.1,5.5,5.6,5.0],
'Play':[True,True,False,True,False,True]}
```

```
In [6]: data
```

```
Out[6]: {'Name': ['Ali', 'Ahmed', 'Sara', 'Tom', 'Mike', 'Maria'],
'Gender': ['Male', 'Male', 'Female', 'Male', 'Male', 'Female'],
'Age': [20, 30, 15, 24, 23, 20],
'Height': [5.2, 5.4, 5.1, 5.5, 5.6, 5.0],
'Play': [True, True, False, True, False, True]}
```

```
In [7]: type(data)
```

```
Out[7]: dict
```

3. Converting a Data Dictionary into a Dataframe

```
In [8]: data=pd.DataFrame(data)
```

```
In [9]: data
```

```
Out[9]:
```

	Name	Gender	Age	Height	Play
0	Ali	Male	20	5.2	True
1	Ahmed	Male	30	5.4	True
2	Sara	Female	15	5.1	False
3	Tom	Male	24	5.5	True
4	Mike	Male	23	5.6	False
5	Maria	Female	20	5.0	True

4. Some Basic Commands to Describe a Dataframe

i) shape()

```
In [10]: data.shape
```

```
Out[10]: (6, 5)
```

There are 6 rows and five columns- There is an index as well

ii) columns

```
In [11]: data.columns
```

```
Out[11]: Index(['Name', 'Gender', 'Age', 'Height', 'Play'], dtype='object')
```

data.columns shows the column names

iii) info()

In [12]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 5 columns):
 #   Column   Non-Null Count  Dtype  
---  --       -----          ----  
 0   Name      6 non-null    object  
 1   Gender    6 non-null    object  
 2   Age       6 non-null    int64  
 3   Height    6 non-null    float64 
 4   Play      6 non-null    bool    
dtypes: bool(1), float64(1), int64(1), object(2)
memory usage: 326.0+ bytes
```

info() command displays a detailed information about the dataset including :

1. RangeIndex (Range of Index)
2. Data Columns (Total number of columns)
3. List of all columns with column index , name , Count of Non Null Values and data type
4. Information about data types
5. Memory usage

iv). head() and tail()

In [13]: data.head()

Out[13]:

	Name	Gender	Age	Height	Play
0	Ali	Male	20	5.2	True
1	Ahmed	Male	30	5.4	True
2	Sara	Female	15	5.1	False
3	Tom	Male	24	5.5	True
4	Mike	Male	23	5.6	False

head() shows 5 records from the top whereas tail() shows 5 records from the tail. we can give a parameter to indicate how many rows are required.

```
In [14]: data.head(3)
```

Out[14]:

	Name	Gender	Age	Height	Play
0	Ali	Male	20	5.2	True
1	Ahmed	Male	30	5.4	True
2	Sara	Female	15	5.1	False

```
In [15]: data.tail(3)
```

Out[15]:

	Name	Gender	Age	Height	Play
3	Tom	Male	24	5.5	True
4	Mike	Male	23	5.6	False
5	Maria	Female	20	5.0	True

v) describe()

```
In [16]: data.describe()
```

Out[16]:

	Age	Height
count	6.000000	6.000000
mean	22.000000	5.300000
std	5.01996	0.236643
min	15.000000	5.000000
25%	20.000000	5.125000
50%	21.500000	5.300000
75%	23.750000	5.475000
max	30.000000	5.600000

describe() command shows some statistical measures including count, mean, standard deviation, minimum, Q1, Q2(Median), Q3 and maximum for all the numeric attributes. By default it does not show other types of attributes. we can include them by defining the include parameter.

```
In [17]: data.describe(include=['object'])
```

Out[17]:

	Name	Gender
count	6	6
unique	6	2
top	Mike	Male
freq	1	4

```
In [18]: data.describe(include=['object', 'integer','bool','float'])
```

Out[18]:

	Name	Gender	Age	Height	Play
count	6	6	6.00000	6.000000	6
unique	6	2	NaN	NaN	2
top	Mike	Male	NaN	NaN	True
freq	1	4	NaN	NaN	4
mean	NaN	NaN	22.00000	5.300000	NaN
std	NaN	NaN	5.01996	0.236643	NaN
min	NaN	NaN	15.00000	5.000000	NaN
25%	NaN	NaN	20.00000	5.125000	NaN
50%	NaN	NaN	21.50000	5.300000	NaN
75%	NaN	NaN	23.75000	5.475000	NaN
max	NaN	NaN	30.00000	5.600000	NaN

vi) dtypes

```
In [19]: data.dtypes
```

```
Out[19]: Name      object
          Gender    object
          Age       int64
          Height   float64
          Play      bool
          dtype: object
```

dtypes shows the data types of all the columns in the dataframe

vii) Some commands to get basic statistical measures for individual columns

```
In [20]: data['Age'].mean()
```

```
Out[20]: 22.0
```

```
In [21]: data['Age'].median()
```

```
Out[21]: 21.5
```

```
In [22]: data['Age'].min()
```

```
Out[22]: 15
```

```
In [23]: data['Age'].max()
```

```
Out[23]: 30
```

```
In [24]: data['Age'].var()
```

```
Out[24]: 25.2
```

```
In [25]: Q1=np.quantile(data['Age'],0.25)
Q1
```

```
Out[25]: 20.0
```

```
In [26]: Q2=np.quantile(data['Age'],0.50)
Q2
```

```
Out[26]: 21.5
```

```
In [27]: Q3=np.quantile(data['Age'],0.75)
Q3
```

```
Out[27]: 23.75
```

5). Reading CSV/ Excel files into a Dataframe

```
In [28]: data=pd.read_csv("D:/FTI/Cohort 2 EDA/Lecture 1/sampleCSV.csv")
```

```
In [29]: data.head()
```

```
Out[29]:
```

S.NO	Name	Age	Gender	
0	1	Ali	20	Male
1	2	Ahmed	21	Male
2	3	Usman	22	Male
3	4	Akram	23	Male
4	5	Tom	19	Male

```
In [30]: data=pd.read_csv("D:/FTI/Cohort 2 EDA/Lecture 1/sampleCSVNoHeader.csv", header=None)
```

```
In [31]: data
```

```
Out[31]:
```

0	1	2	3	
0	1	Ali	20	Male
1	2	Ahmed	21	Male
2	3	Usman	22	Male
3	4	Akram	23	Male
4	5	Tom	19	Male
5	6	Sara	17	Female
6	7	David	23	Male
7	8	Dravid	22	Male
8	9	Bob	18	Male
9	10	Clara	21	Female

```
In [32]: data=pd.read_excel("D:/FTI/Cohort 2 EDA/Lecture 1/sampleExcel.xlsx")
```

```
In [33]: data.head()
```

Out[33]:

S.NO	Name	Age	Gender
0	1	Ali	20
1	2	Ahmed	21
2	3	Usman	22
3	4	Akram	23
4	5	Tom	19

```
In [34]: data=pd.read_excel("D:/FTI/Cohort 2 EDA/Lecture 1/sampleExcelNoHeader.xlsx", header=None)
```

```
In [35]: data.head()
```

Out[35]:

0	1	2	3
0	1	Ali	20
1	2	Ahmed	21
2	3	Usman	22
3	4	Akram	23
4	5	Tom	19

```
In [36]: data.columns=['ID', 'Name', 'Age', 'Gender']
```

```
In [37]: data.head()
```

Out[37]:

ID	Name	Age	Gender
0	1	Ali	20
1	2	Ahmed	21
2	3	Usman	22
3	4	Akram	23
4	5	Tom	19

6. Saving a dataframe as a csv

```
In [38]: data.to_csv("D:/FTI/Cohort 2 EDA/Lecture 1/data.csv", index=False)
```

7. Working with Telecom Churn Dataset

i) Load the dataset

```
In [39]: data=pd.read_csv("D:/FTI/Cohort 2 EDA/Lecture 1/TelecomChurn.csv")
```

In [40]: `data.head()`

Out[40]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	total eve calls	total eve charge	...
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	99	16.78	...
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	103	16.62	...
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	110	10.30	...
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	88	5.26	...
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	122	12.61	...

5 rows × 21 columns

ii) Basis Data Description

In [41]: `data.shape`

Out[41]: (3333, 21)

In [42]: `pd.set_option('Display.Max_Columns', 50)`

In [43]: `data.head()`

Out[43]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	total eve minutes	total eve calls	total eve charge	...
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	197.4	99	16.78	16.78	
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	195.5	103	16.62	16.62	
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	121.2	110	10.30	10.30	
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	61.9	88	5.26	5.26	
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	148.3	122	12.61	12.61	

In [44]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   state            3333 non-null    object  
 1   account length   3333 non-null    int64  
 2   area code         3333 non-null    int64  
 3   phone number     3333 non-null    object  
 4   international plan 3333 non-null    object  
 5   voice mail plan  3333 non-null    object  
 6   number vmail messages 3333 non-null    int64  
 7   total day minutes 3333 non-null    float64 
 8   total day calls   3333 non-null    int64  
 9   total day charge  3333 non-null    float64 
 10  total eve minutes 3333 non-null    float64 
 11  total eve calls   3333 non-null    int64  
 12  total eve charge  3333 non-null    float64 
 13  total night minutes 3333 non-null    float64 
 14  total night calls  3333 non-null    int64  
 15  total night charge 3333 non-null    float64 
 16  total intl minutes 3333 non-null    float64 
 17  total intl calls   3333 non-null    int64  
 18  total intl charge  3333 non-null    float64 
 19  customer service calls 3333 non-null    int64  
 20  churn             3333 non-null    bool  
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

In [45]: `data.dtypes`

```
Out[45]: state          object  
account length   int64  
area code         int64  
phone number     object  
international plan  object  
voice mail plan  object  
number vmail messages  int64  
total day minutes float64 
total day calls   int64  
total day charge  float64 
total eve minutes float64 
total eve calls   int64  
total eve charge  float64 
total night minutes float64 
total night calls  int64  
total night charge float64 
total intl minutes float64 
total intl calls   int64  
total intl charge  float64 
customer service calls  int64  
churn             bool  
dtype: object
```

In [46]: `data.describe()`

Out[46]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total ev call
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.980348	100.11431
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.713844	19.92262
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000	87.000000
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000	100.000000
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000	114.000000
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000	170.000000

In [47]: `data.describe(include=['object'])`

Out[47]:

	state	phone number	international plan	voice mail plan
count	3333	3333	3333	3333
unique	51	3333	2	2
top	WV	381-7597	no	no
freq	106	1	3010	2411

In [48]: `data.describe(include=['bool'])`

Out[48]:

	churn
count	3333
unique	2
top	False
freq	2850

In [49]: `data.head(10)`

Out[49]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	197.4	99	16
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	195.5	103	16
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	121.2	110	10
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	61.9	88	5
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	148.3	122	12
5	AL	118	510	391-8027	yes	no	0	223.4	98	37.98	220.6	101	18
6	MA	121	510	355-9993	no	yes	24	218.2	88	37.09	348.5	108	28
7	MO	147	415	329-9001	yes	no	0	157.0	79	26.69	103.1	94	8
8	LA	117	408	335-4719	no	no	0	184.5	97	31.37	351.6	80	28
9	WV	141	415	330-8173	yes	yes	37	258.6	84	43.96	222.0	111	18

In [50]: `data.columns`

Out[50]: `Index(['state', 'account length', 'area code', 'phone number', 'international plan', 'voice mail plan', 'number vmail messages', 'total day minutes', 'total day calls', 'total day charge', 'total eve minutes', 'total eve calls', 'total eve charge', 'total night minutes', 'total night calls', 'total night charge', 'total intl minutes', 'total intl calls', 'total intl charge', 'customer service calls', 'churn'], dtype='object')`

8) Univariate Analysis

i) Categorical and Boolean Attributes

In [51]: `data.describe(include=['object', 'bool'])`

Out[51]:

	state	phone number	international plan	voice mail plan	churn
count	3333	3333	3333	3333	3333
unique	51	3333	2	2	2
top	WV	381-7597	no	no	False
freq	106	1	3010	2411	2850

a) `describe()`

```
In [52]: data['state'].describe()
```

```
Out[52]: count    3333
unique     51
top       WV
freq     106
Name: state, dtype: object
```

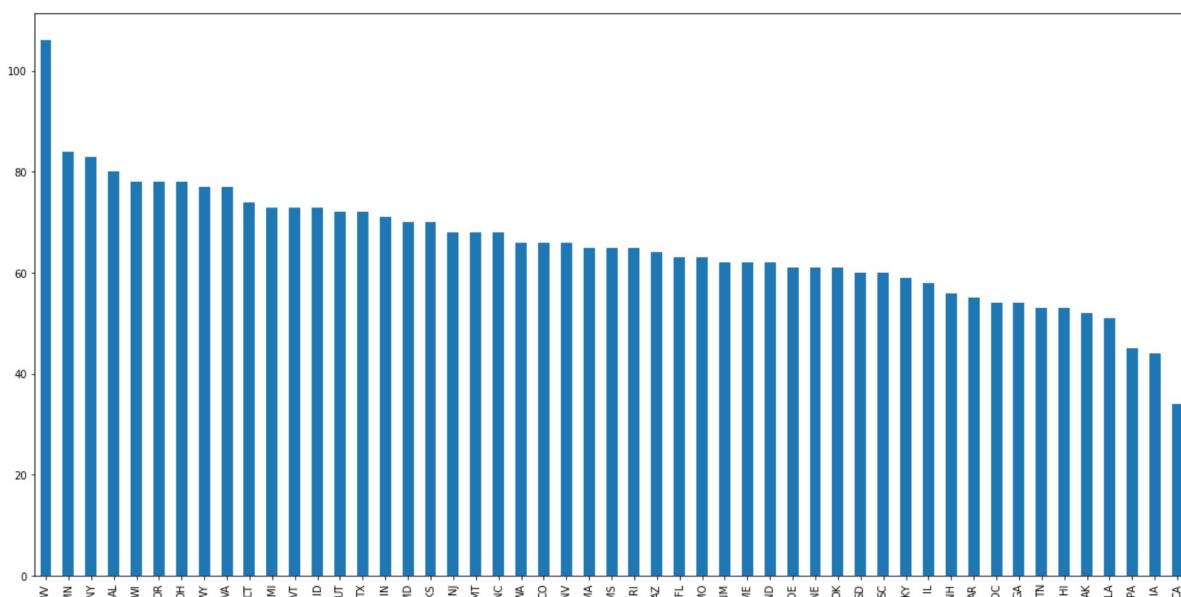
b) value_counts

```
In [53]: data['state'].value_counts();
```

c) Bar Chart

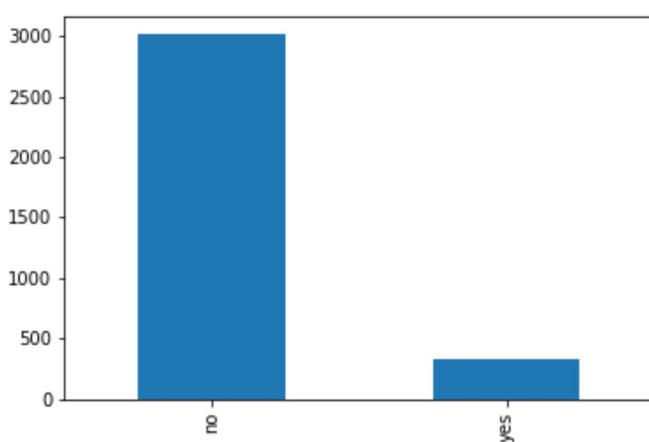
```
In [171]: data['state'].value_counts().plot(kind='bar', figsize=(20,10))
```

```
Out[171]: <matplotlib.axes._subplots.AxesSubplot at 0x1be37a13f28>
```



```
In [170]: data['international plan'].value_counts().plot(kind='bar')
```

```
Out[170]: <matplotlib.axes._subplots.AxesSubplot at 0x1be379c1748>
```



ii) Numeric Attributes

In [172]: `data.describe()`

Out[172]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total ev call
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.980348	100.11431
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.713844	19.92262
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000	87.000000
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000	100.000000
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000	114.000000
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000	170.000000

a) describe()

In [173]: `data['total day charge'].describe()`

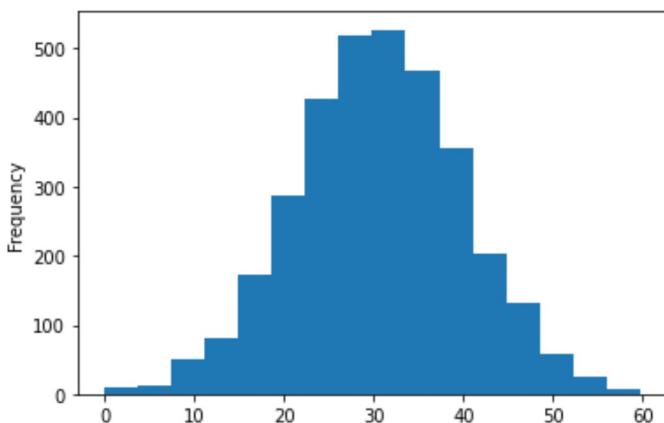
Out[173]:

```
count      3333.000000
mean       30.562307
std        9.259435
min        0.000000
25%       24.430000
50%       30.500000
75%       36.790000
max       59.640000
Name: total day charge, dtype: float64
```

b) histogram

In [188]: `data['total day charge'].plot.hist(bins=16)`

Out[188]: <matplotlib.axes._subplots.AxesSubplot at 0x1be373b6828>



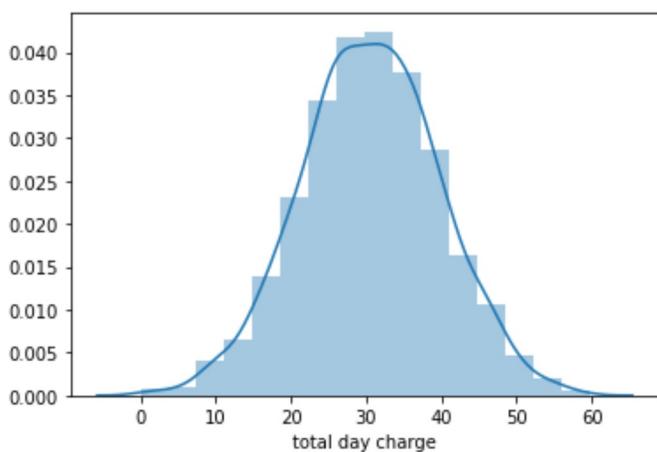
```
In [54]: import seaborn as sns
```

```
In [55]: sns.distplot(data['total day charge'], bins=16, kde=True)
```

C:\Users\malik\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

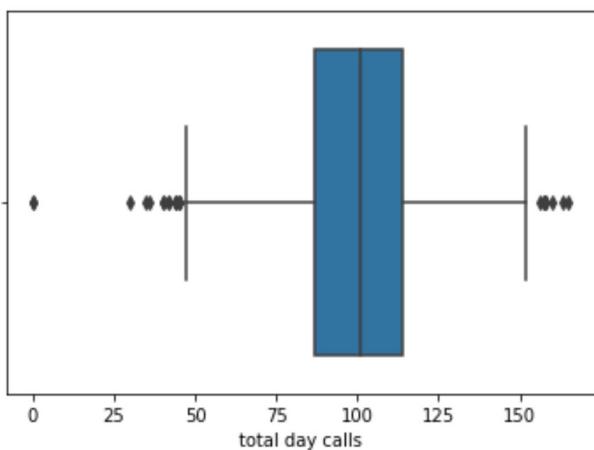
```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x22057c5a668>
```



c) boxplot

```
In [193]: sns.boxplot(x='total day calls', data=data)
```

```
Out[193]: <matplotlib.axes._subplots.AxesSubplot at 0x1be3ab3c320>
```



9. Bivariate Analysis

```
In [194]: data['churn'].value_counts()
```

```
Out[194]: False      2850
          True       483
          Name: churn, dtype: int64
```

In [195]: `data.head()`

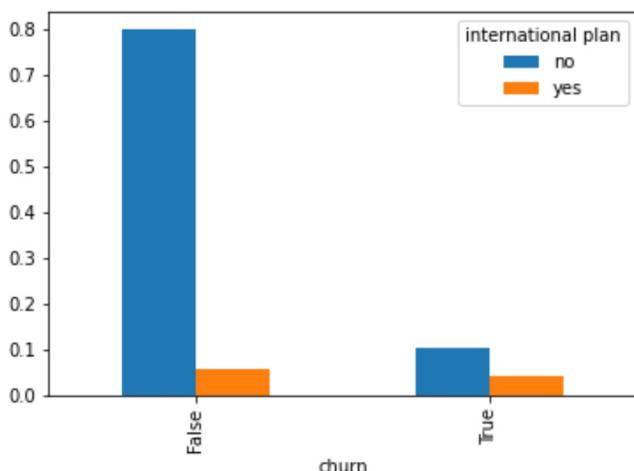
Out[195]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total charge
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	197.4	99	16
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	195.5	103	16
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	121.2	110	10
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	61.9	88	5
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	148.3	122	12

a) cross tab [Categorical vs Categorical]

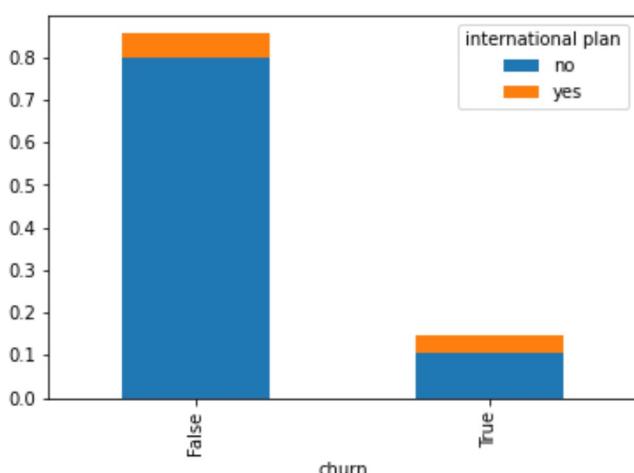
In [199]: `pd.crosstab(data['churn'], data['international plan'], normalize=True).plot(kind='bar')`

Out[199]: <matplotlib.axes._subplots.AxesSubplot at 0x1be3ab9c908>



In [200]: `pd.crosstab(data['churn'], data['international plan'], normalize=True).plot(kind='bar', stacked=True)`

Out[200]: <matplotlib.axes._subplots.AxesSubplot at 0x1be3acb5cc0>



b) Categorical vs Numeric (groupby)

Whether churners make more customer calls than non-churners

```
In [212]: data.groupby('churn')['customer service calls'].mean()
```

```
Out[212]: churn
False    1.449825
True     2.229814
Name: customer service calls, dtype: float64
```

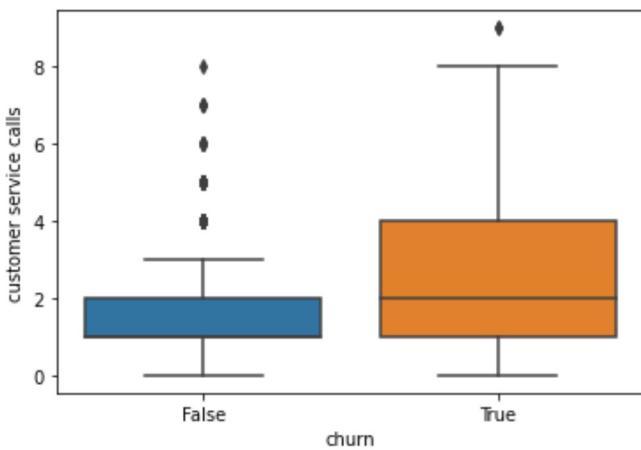
which state has the highest number of churning

```
In [257]: data.groupby('state')['churn'].value_counts().head(20)
```

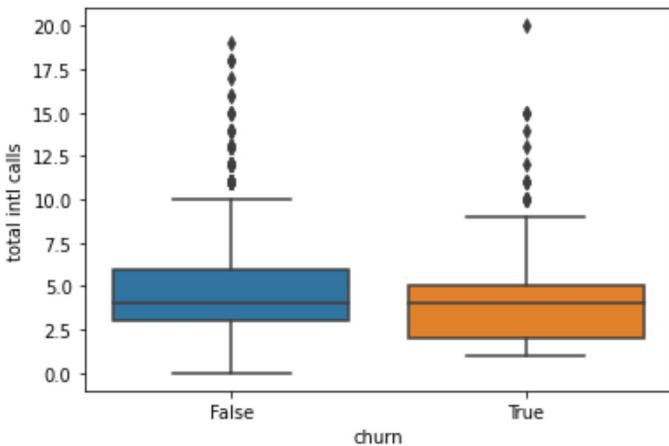
```
Out[257]: state  churn
AK      False    49
          True     3
AL      False    72
          True     8
AR      False    44
          True    11
AZ      False    60
          True     4
CA      False    25
          True     9
CO      False    57
          True     9
CT      False    62
          True    12
DC      False    49
          True     5
DE      False    52
          True     9
FL      False    55
          True     8
Name: churn, dtype: int64
```

c) boxplots

```
In [273]: sns.boxplot( x=data["churn"], y=data["customer service calls"] );
plt.show()
```



```
In [275]: sns.boxplot( x=data["churn"], y=data["total intl calls"] );
plt.show()
```



```
In [ ]:
```