

EDA 04- Feature Selection -I

1. Necessary Imports

```
In [25]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
import scipy.stats as s
import seaborn as sns
```

2. Reading dataset into CSV & Basic Data Description

a) Reading Data

```
In [87]: data=pd.read_csv("D:/FTI/Cohort 2 EDA/Lecture 4/Finance.csv")
```

```
In [4]: data.head()
```

Out[4]:

	age	job	marital	education	default	balance	housing	loan	contact	day	n
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	
1	44	technician	single	secondary	no	29	yes	no	unknown	5	
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	
4	33	unknown	single	unknown	no	1	no	no	unknown	5	

```
In [5]: data.shape
```

Out[5]: (45211, 17)

b) Check the Data Types

```
In [79]: data.dtypes
```

```
Out[79]: age           int64
job           object
marital       object
education     object
default       object
balance       int64
housing       object
loan          object
contact       object
day           int64
month         object
duration      int64
campaign      int64
pdays        int64
previous      int64
poutcome     object
y            object
dtype: object
```

3. Check Missing Data or Null Values

```
In [109]: data.isnull().sum()
```

```
Out[109]: age           0
job           0
marital       0
education     0
default       0
balance       0
housing       0
loan          0
contact       0
day           0
month         0
duration      0
campaign      0
pdays        0
previous      0
poutcome     0
y            0
dtype: int64
```

4. Check for Data Quality Issues

a) For categorical attributes you can inspect the unique values

You can display all the unique values and based on domain knowledge can decide if incorrect data exists or not.

```
In [10]: data['job'].unique()

Out[10]: array(['management', 'technician', 'entrepreneur', 'blue-collar',
                'unknown', 'retired', 'admin.', 'services', 'self-employed',
                'unemployed', 'housemaid', 'student'], dtype=object)
```

If you do not have the domain knowledge, then value counts may give you an idea about the possible incorrect values

```
In [11]: data['job'].value_counts()

Out[11]: blue-collar      9732
         management      9458
         technician      7597
         admin.          5171
         services        4154
         retired         2264
         self-employed    1579
         entrepreneur    1487
         unemployed      1303
         housemaid       1240
         student         938
         unknown         288
         Name: job, dtype: int64
```

You can also use describe() to get a bit more information about the feature under consideration

```
In [12]: data['job'].describe()

Out[12]: count      45211
         unique        12
         top    blue-collar
         freq      9732
         Name: job, dtype: object
```

b) For numeric attributes you can inspect all the possible set of values

If you have a domain knowledge, the minimum and maximum values can spot if incorrect data is present or not


```
In [14]: data['age'].describe()
```

```
Out[14]: count      45211.000000  
mean         40.936210  
std          10.618762  
min          18.000000  
25%          33.000000  
50%          39.000000  
75%          48.000000  
max          95.000000  
Name: age, dtype: float64
```

If you do not have the domain knowledge, the boxplot may give you some clue about the presence of possibly incorrect values. We call them as outliers. Outliers may be real or due to data collection problems.

```
In [19]: sns.boxplot(x=data['age'], data=data)
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x1e8e9b634a8>
```



5. Feature Selection Based on Filter Methods

a) Using Chi Square to test association between categorical attributes

The class variable is Y which is object type. So we can use Chi Square to check the association between all other object types and the class. The variables having strong association can be chosen as the features for machine learning algorithm.

Chi-Square Hypothesis : HO: There is no relationship / association / dependence between two attributes H1: There is a relationship / associatio/ dependence between two variables

Lets perform a chi-square test of independence for 'default' and class variable 'y'

```
In [20]: data['default'].value_counts()
```

```
Out[20]: no      44396
         yes      815
         Name: default, dtype: int64
```

```
In [21]: data['y'].value_counts()
```

```
Out[21]: no      39922
         yes      5289
         Name: y, dtype: int64
```

```
In [22]: ov=pd.crosstab(data['default'],data['y'])
```

```
In [23]: ov
```

```
Out[23]:
```

	y	no	yes
default			
no	39159	5237	
yes	763	52	

```
In [26]: b=s.chi2_contingency(ov)
```

```
In [27]: b
```

```
Out[27]: (22.20224995571685,
          2.4538606753508344e-06,
          1,
          array([[39202.34261574,  5193.65738426],
                 [ 719.65738426,    95.34261574]]))
```

Here b is a tuple containing an immutable sequence of python objects . Here it contains four objects . b[0] contains the value of chi2 statistic , b[1] contains the p-value of the test , b[2] contains the degree of freedom and b[3] contains the expected frequencies.

```
In [119]: b[0]
```

```
Out[119]: 22.20224995571685
```

```
In [120]: b[1]
```

```
Out[120]: 2.4538606753508344e-06
```

```
In [121]: b[2]
```

```
Out[121]: 1
```

```
In [122]: b[3]
```

```
Out[122]: array([[39202.34261574,  5193.65738426],
                 [ 719.65738426,    95.34261574]])
```

Lets create a custom function to perform chi-square test of independence

```
In [123]: def test_dependency(data, f1, f2, alpha):
            ov=pd.crosstab(data[f1],data[f2])
            b=s.chi2_contingency(ov)
            chi2_statistic=b[0]
            p_value=b[1]
            dof=b[2]
            critical_value=s.chi2.ppf(q=1-alpha, df=dof)
            print('Significance level: ',alpha)
            print('Degree of Freedom: ',dof)
            print('chi-square statistic:',chi2_statistic)
            print('critical_value:',critical_value)
            print('p-value:',p_value)

            if chi2_statistic>=critical_value:
                print("Reject H0,There is a relationship between 2 categorical variables")
            else:
                print("Retain H0,There is no relationship between 2 categorical variables")

            if p_value<=alpha:
                print("Reject H0,There is a relationship between 2 categorical variables")
            else:
                print("Retain H0,There is no relationship between 2 categorical variables")
```

```
In [126]: test_dependency(data, 'default', 'y', 0.05)
```

```
Significance level:  0.05
Degree of Freedom:  1
chi-square statistic: 22.20224995571685
critical_value: 3.841458820694124
p-value: 2.4538606753508344e-06
Reject H0,There is a relationship between 2 categorical variables
Reject H0,There is a relationship between 2 categorical variables
```

```
In [125]: test_dependency(data, 'education', 'y', 0.05)
```

```
Significance level:  0.05
Degree of Freedom:  3
chi-square statistic: 238.92350616407606
critical_value: 7.814727903251179
p-value: 1.6266562124072994e-51
Reject H0,There is a relationship between 2 categorical variables
Reject H0,There is a relationship between 2 categorical variables
```

Selecting k-Best Features based on Chi-Square Test

We will be using `SelectKBest()` which takes numeric data only. So for that we have to encode all the categorical. We will be using manual encoding for ordinal features whereas label encoding for all other nominal features.

Encoding Ordinal Features

```
In [59]: ordinal_list=['education']
data['education'] = data['education'].replace(['primary', 'secondary', 'tertiary', 'unknown'], [1, 2, 3, 2])
```

```
In [60]: data.head()
```

Out[60]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	du
0	58	4	1	3	0	2143	1	0	2	5	8	
1	44	9	2	2	0	29	1	0	2	5	8	
2	33	2	1	2	0	2	1	1	2	5	8	
3	47	1	1	2	0	1506	1	0	2	5	8	
4	33	11	2	2	0	1	0	0	2	5	8	

Encoding Nominal Features

```
In [56]: nominal_list = []
for i in data.columns.tolist():
    if (data[i].dtype=='object') and (i not in ordinal_list):
        nominal_list.append(i)
print (nominal_list)
print ('Number of nominal features:', str(len(nominal_list)))

['job', 'marital', 'default', 'housing', 'loan', 'contact', 'month', 'poutcome', 'y']
Number of nominal features: 9
```

```
In [57]: from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
for column in nominal_list:
    data[column]=encoder.fit_transform(data[column])
```

```
In [61]: data.head()
```

```
Out[61]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	du
0	58	4	1	3	0	2143	1	0	2	5	8	
1	44	9	2	2	0	29	1	0	2	5	8	
2	33	2	1	2	0	2	1	1	2	5	8	
3	47	1	1	2	0	1506	1	0	2	5	8	
4	33	11	2	2	0	1	0	0	2	5	8	

Combining ordinal and nominal features after encoding

```
In [73]: ordinal_data=data[ordinal_list]
nominal_data=data[nominal_list]
categorical_data = pd.concat([ordinal_data,nominal_data], axis=1)
```

```
In [72]: categorical_data
```

```
Out[72]:
```

	job	marital	default	housing	loan	contact	month	poutcome	y
0	4	1	0	1	0	2	8	3	0
1	9	2	0	1	0	2	8	3	0
2	2	1	0	1	1	2	8	3	0
3	1	1	0	1	0	2	8	3	0
4	11	2	0	0	0	2	8	3	0
...
45206	9	1	0	0	0	0	9	3	1
45207	5	0	0	0	0	0	9	3	1
45208	5	1	0	0	0	0	9	2	1
45209	1	1	0	0	0	1	9	3	0
45210	2	1	0	0	0	0	9	1	0

45211 rows × 9 columns

Selecting K Best Features based on Chi-Square Test


```
In [74]: X=categorical_data.drop('y',axis=1)
Y=categorical_data['y']
chi_scores = chi2(X,Y)
```

```
In [75]: p_values = pd.Series(chi_scores[1],index = X.columns)
p_values.sort_values(ascending = True , inplace = True)
chi2_values=pd.Series(chi_scores[0],index = X.columns)
chi2_values.sort_values(ascending = False , inplace = True)
```




```
In [76]: p_values.plot.bar()
```

```
Out[76]: <matplotlib.axes._subplots.AxesSubplot at 0x1e8ea816748>
```



```
In [77]: chi2_values.plot.bar()
```

```
Out[77]: <matplotlib.axes._subplots.AxesSubplot at 0x1e8eb4231d0>
```



```
In [78]: # Three features with highest chi-squared statistics are selected
chi2_features = SelectKBest(chi2, k = 3)
X_kbest_features = chi2_features.fit_transform(X, Y)

# Reduced features
print('Original feature number:', X.shape[1])
print('Reduced feature number:', X_kbest_features.shape[1])
```

```
Original feature number: 9
```

```
Reduced feature number: 3
```

```
In [81]: index = chi2_features.get_support(indices=True)
         print(index)
```

```
[1 4 6]
```

```
In [82]: X
```

```
Out[82]:
```

	education	job	marital	default	housing	loan	contact	month	poutcome
0	3	4	1	0	1	0	2	8	3
1	2	9	2	0	1	0	2	8	3
2	2	2	1	0	1	1	2	8	3
3	2	1	1	0	1	0	2	8	3
4	2	11	2	0	0	0	2	8	3
...
45206	3	9	1	0	0	0	0	9	3
45207	1	5	0	0	0	0	0	9	3
45208	2	5	1	0	0	0	0	9	2
45209	2	1	1	0	0	0	1	9	3
45210	2	2	1	0	0	0	0	9	1

```
45211 rows × 9 columns
```

Features at index 1 , 4 and 6 are job , housing and contact respectively

```
In [85]: selected_features= ['job','housing','contact']
```

b) Using Pearson Correlation Coefficient for Numeric Features vs Numeric Class

```
In [136]: data=pd.read_csv("D:/FTI/Cohort 2 EDA/Lecture 4/Finance.csv")
```

```
In [137]: df=data.copy()
```

```
In [138]: data['y'].dtype
```

```
Out[138]: dtype('O')
```

```
In [139]: data['y'].value_counts()
```

```
Out[139]: no      39922
         yes      5289
         Name: y, dtype: int64
```

we can convert yes and no into 1 and 0 and change the data type from object to integer

```
In [140]: data['y']=data['y'].apply(lambda x:0 if x=='no' else 1)
```

```
In [141]: data['y'].value_counts()
```

```
Out[141]: 0    39922
          1     5289
          Name: y, dtype: int64
```

```
In [142]: data['y'].dtype
```

```
Out[142]: dtype('int64')
```

```
In [143]: cor_matrix=data.corr()
```

```
In [144]: print(cor_matrix['y'].sort_values(ascending=False))
```

```
y          1.000000
duration    0.394521
pdays      0.103621
previous    0.093236
balance     0.052838
age         0.025155
day        -0.028348
campaign    -0.073172
Name: y, dtype: float64
```

```
In [145]: cor_matrix
```

```
Out[145]:
```

	age	balance	day	duration	campaign	pdays	previous	y
age	1.000000	0.097783	-0.009120	-0.004648	0.004760	-0.023758	0.001288	0.025155
balance	0.097783	1.000000	0.004503	0.021560	-0.014578	0.003435	0.016674	0.052838
day	-0.009120	0.004503	1.000000	-0.030206	0.162490	-0.093044	-0.051710	-0.028348
duration	-0.004648	0.021560	-0.030206	1.000000	-0.084570	-0.001565	0.001203	0.394521
campaign	0.004760	-0.014578	0.162490	-0.084570	1.000000	-0.088628	-0.032855	-0.073172
pdays	-0.023758	0.003435	-0.093044	-0.001565	-0.088628	1.000000	0.454820	0.103621
previous	0.001288	0.016674	-0.051710	0.001203	-0.032855	0.454820	1.000000	0.093236
y	0.025155	0.052838	-0.028348	0.394521	-0.073172	0.103621	0.093236	1.000000

```
In [146]: features_cor=(cor_matrix['y'].sort_values(ascending=False))
```

```
In [147]: selected_num_features= ['duration','pdays','previous']
```

```
In [148]: best_features=selected_features+(selected_num_features)
```

```
In [149]: best_features
```

```
Out[149]: ['job', 'housing', 'contact', 'duration', 'pdays', 'previous']
```

```
In [151]: data2=pd.concat([data[best_features],data['y']],axis=1)
```

In [152]: data2

Out[152]:

	job	housing	contact	duration	pdays	previous	y
0	management	yes	unknown	261	-1	0	0
1	technician	yes	unknown	151	-1	0	0
2	entrepreneur	yes	unknown	76	-1	0	0
3	blue-collar	yes	unknown	92	-1	0	0
4	unknown	no	unknown	198	-1	0	0
...
45206	technician	no	cellular	977	-1	0	1
45207	retired	no	cellular	456	-1	0	1
45208	retired	no	cellular	1127	184	3	1
45209	blue-collar	no	telephone	508	-1	0	0
45210	entrepreneur	no	cellular	361	188	11	0

45211 rows × 7 columns

Remarks: Feature Selection using filter methods is independant of the machine learning model we use.
Whether our selected features will work better or not depends on the performance of ML model we apply.